

Welcome to A Catholic Introduction to Artificial Intelligence Class 6! Ethics of AI – Bias & Fairness

"AI is a witness to truth at best, never the source of truth."

Our guide for this class is *Magnifica Humanitas* ("Magnificent Humanity"), Pope Leo XIV's first encyclical, *On Safeguarding the Human Person in the Time of Artificial Intelligence*, signed on May 15, 2026.

Learning Objectives

By the end of this class, you will learn about how AI can reflect biases found in its training data, leading to discrimination in hiring, policing, and facial recognition. Case studies will highlight real-world examples of AI bias and its consequences. The class will explore ways to make AI more ethical and fairer, followed by a discussion on who is responsible for AI's actions.

Specifically, by the end of this class you should be able to:

1. Explain, using *Magnifica Humanitas*, why AI imitates certain functions of human intelligence without understanding what it produces or knowing truth.
2. Explain why Pope Leo XIV teaches that technology is "never neutral."
3. Describe how an AI system's "impression of objectivity" can hide the biases of those who design and train it, and why this matters for justice.
4. Explain what the Pope means by "truth as a common good" and by "power detached from truth."

5. Distinguish information from wisdom, and explain why learning to use AI well includes learning when *not* to use it.
6. Apply the Church's teaching on human dignity and conscience to the moral limits of what we hand over to machines.

1. Class 5 Review

In Class 5, we explored one of the most important technologies in modern artificial intelligence: deep learning. Deep learning is a branch of machine learning that uses layered neural networks to recognize patterns and make decisions from large amounts of data. We learned that artificial neural networks are loosely inspired by the structure of the human brain, although they are mathematical systems rather than actual digital brains. These networks contain layers of interconnected “neurons” that process information step-by-step, allowing computers to identify increasingly complex features in images, sound, and language.

The class also examined how neural networks learn through training. We were introduced to the idea of backpropagation, where a model measures its mistakes and adjusts itself to improve future predictions. This process allows deep learning systems to become highly effective at tasks such as image recognition, speech recognition, translation, and text generation. Students discussed how modern AI systems rely on enormous amounts of data and computing power in order to achieve strong performance.

A major focus of the lesson was Large Language Models (LLMs). We learned that LLMs generate language by predicting the next token based on patterns learned during training. The class explored where these models obtain their information, including books, articles, websites, and other large collections of digital text. Students learned that LLMs do not “think” or “understand” like humans, but instead recognize statistical relationships in language. This helped clarify why AI systems can often produce fluent and convincing responses even when they are incorrect.

We also discussed several common misconceptions about AI. These included the belief that AI systems truly understand meaning like humans, that ChatGPT searches the internet for every answer, or that AI systems always know when they are wrong. The lesson emphasized that deep learning systems are powerful pattern-recognition tools, but they can still produce hallucinations, make mistakes, and fail in unfamiliar situations. Students were encouraged to think critically about AI-generated information rather than assuming confident answers are automatically correct.

2. Limitations of LLMs

- **Lack of Up-to-Date Information**: Training LLMs is an expensive and time-intensive task. It takes hundreds of GPUs running on hundreds of servers. Few companies can

afford to do this. Also, because of these constraints, a model's knowledge is capped at a certain date. Models can be two to three years behind the current time; its `knowledge` is `static` - locked in time. Its `knowledge` is a statistical reflection of patterns in their training data. They will confidently state incorrect or outdated facts (a problem called hallucination). It will never tell you it made it up.

We will learn about how these limitations are overcome is subsequent classes.

- **Hallucinations**: When a model doesn't know an answer, it will make one up! Without any safeguards to prevent this from happening and also because models are solely statistical, it will pull from what it knows (has been `trained` on) and produce an answer without even knowing if it is true or not.
- **Lack of understanding of the Truth, Bias, Ethics, and Morality**: The degree to which any LLM is truthful, moral, or ethical is solely proportional to the data that was used during the training process.
- **Not fully understood, even by its makers**: Pope Leo XIV observes that today's AI systems are more "cultivated" than "built." Developers build a framework and the model "grows" from data, so even they do not fully know how it arrives at a particular answer. That is one more reason a confident AI answer must still be checked.

> "Current AI systems are more “cultivated” than “built,” for developers do not directly design every detail, but instead create a framework within which the intelligence “grows.”"
>
> — *Magnifica Humanitas*, no. 98

3. *Magnifica Humanitas*: Babel or Jerusalem?

On May 15, 2026, the 135th anniversary of Pope Leo XIII's *Rerum Novarum*, Pope Leo XIV signed *Magnifica Humanitas*, the first encyclical in the Church's history devoted to artificial intelligence. Leo XIII applied the Gospel to the "new things" of the industrial revolution. Leo XIV applies the Church's social teaching to the new things of our own time: digital platforms, algorithms, and AI.

The encyclical opens with a choice:

> "Humanity, created by God in all its grandeur, is today facing a pivotal choice: either to construct a new Tower of Babel or to build the city in which God and humanity dwell together."
>
> — *Magnifica Humanitas*, no. 1

The Pope uses two scenes from Scripture to describe that choice:

- **The Tower of Babel** (Genesis 11): a project "conceived without reference to God," with "a single language, a single technology, a single direction." It is built on pride, chooses "homogenization over communion," and "sacrifices human dignity for efficiency."
- **Nehemiah rebuilding the walls of Jerusalem** (Nehemiah 2–6): before acting, Nehemiah "fasted, prayed and interceded for the people." He did not impose solutions from above. He gave each family its own section of the wall to rebuild. It is "an undertaking with God at the center, which rebuilds relationships before rebuilding with stones."

So the question is not whether to accept or reject AI. The Pope says the real choice is "between constructing Babel or rebuilding Jerusalem." And he gives the reason in one sentence that frames this entire class:

> "In the abstract, technology in and of itself is not a solution to humanity's problems, just as it is not inherently evil. In practice, however, technology is never neutral, because it takes on the characteristics of those who devise, finance, regulate and use it."
>
> — *Magnifica Humanitas*, no. 9

If technology takes on the characteristics of the people who make and use it, then an AI system's relationship to **truth**, its **biases**, and the **ethics** and **morality** of how it is used all lead back to human beings and their choices. That is why the Pope ends his chapter on AI by pointing to the human heart. Saint Augustine taught that two loves build two cities: love of God and neighbor, or love of self. The Pope applies this directly to our time:

> "The age of AI is no exception: the construction of Babel or the rebuilding of Jerusalem begins within each one of us."
>
> — *Magnifica Humanitas*, no. 130

The Pope also warns against two opposite mistakes: "We cannot condone naïve enthusiasms, nor fuel unfounded fears." Instead, he asks for discernment, and he names "the promotion of digital literacy" as one concrete way to build the common good. That is what this course is for. His test for any technology comes from Saint John Paul II: does it "make human life on earth 'more human' in every aspect of that life? Does it make it more worthy of man?"

4. Truth

From a Catholic perspective, AI does not "know" truth in the way a human person can know truth. That distinction is fundamental. Human beings are created in the image of God and possess intellect and free will (Genesis 1:27). We are capable of perceiving reality, reasoning morally, and orienting ourselves toward Truth Himself, who is God. AI, however sophisticated, is not a rational soul. It is a tool that processes symbols, probabilities, and

patterns through mathematical systems. It can manipulate statements about truth, but it cannot possess truth interiorly.

Classical Catholic philosophy, especially in the thought of Saint Thomas Aquinas, teaches that truth is the conformity of the intellect to reality — *adaequatio intellectus et rei*. In other words, truth exists when the mind correctly apprehends what is real. AI does not have intellect in the Thomistic sense. It has no soul, no consciousness, no spiritual faculty of reason, no conscience, and no moral agency. It cannot believe, doubt, repent, or assent to truth. It simply generates outputs based on patterns within the data on which it was trained.

Magnifica Humanitas teaches the same thing. The Pope warns against "the misconception of equating this type of 'intelligence' with that of human beings," because these systems "merely imitate certain functions of human intelligence." They are often far faster than we are, "yet this power remains entirely tied to data processing."

> "They may imitate language, behavior and analytical skills, or even simulate empathy and understanding, but they do not understand what they produce, for they lack the affective, relational and spiritual perspective through which human beings grow in wisdom."
>
> — *Magnifica Humanitas*, no. 99

This means AI does not understand truth metaphysically or spiritually. It can produce statements that are true, false, contradictory, misleading, or profound without any awareness of what those statements mean. An AI may accurately explain the Resurrection of Christ in one response and reduce it to symbolism in another, not because it is wrestling with theology, but because it is predicting language patterns from its training and instructions. Only persons can genuinely seek truth, love truth, or reject truth.

Even when we say a model "learns," the Pope reminds us that this is not learning in the human sense. A person grows "through choices, mistakes, forgiveness and fidelity." A model undergoes "a form of statistical adaptation based on data and feedback, which can be very effective, but does not imply inner growth." You saw this in Class 5: training and backpropagation adjust numbers, but they do not make a model wiser.

An AI model predicts likely sequences of words based on data. It can produce statements that are true, false, mixed, deceptive, or incoherent without any interior awareness.

If an AI says:

> “Jesus rose from the dead”

and then elsewhere says:

> “The Resurrection is merely symbolic,”

it is not wrestling with contradiction the way a human theologian would. It is generating outputs from patterns in training data and instructions.

In Catholic terms, AI lacks:

- intellect (*intellectus*)
- reason as a spiritual faculty
- free will
- moral culpability
- conscience
- the capacity for faith

It cannot assent to truth. It cannot commit heresy. It cannot believe. Only persons can do that.

Catholicism also insists that truth is not merely the accumulation of information or consensus opinion. Modern society often treats truth as something determined by statistics, cultural agreement, emotional resonance, or institutional authority. Christianity rejects this reduction. Christ declares in Scripture, “I am the way, the truth, and the life” (John 14:6). Truth is ultimately personal because it is grounded in the very nature of God. Objective reality exists independently of algorithms, governments, corporations, or public opinion.

Pope Leo XIV adds an important insight: **truth is a common good**. Like clean water, truth belongs to everyone, and everyone is harmed when it is polluted. He notes that "disinformation did not begin with AI, yet today it finds a powerful amplifier in AI." He also explains how truth is actually found. It has "a rational dimension, as it requires verification, cross-checking of sources and responsible argumentation," and it is "deeply relational, built through bonds of trust." Truth does not come from a machine or a central authority: "truthful information does not arise from centralized or automated control."

> "Our first task is neither to demonize nor idolize technological tools, but to utilize them on the basis of a fundamental principle, namely that truth is a common good and not the property of those with power or influence."
>
> — *Magnifica Humanitas*, no. 137

Because AI systems are trained on human-produced material, they inevitably reflect the strengths and weaknesses of fallen human culture. Human civilization contains wisdom, beauty, and insight, but also confusion, propaganda, pride, ideology, and moral corruption. AI therefore mirrors both truth and error. A Catholic should never approach AI as an oracle or infallible authority. The machine reflects the disorder and fragmentation of the world from which its data is drawn.

One danger is that AI can simulate authority without possessing wisdom. Human beings naturally associate articulate speech with intelligence and reliability. Yet Catholic tradition distinguishes sharply between information and wisdom. Wisdom involves right judgment,

moral formation, humility before God, and orientation toward the good. AI can summarize theology, philosophy, or Scripture while having no understanding whatsoever of holiness, sacrifice, grace, sin, or salvation. This creates a temptation to replace contemplation with instant answers and discernment with automation.

Magnifica Humanitas names this danger directly. The Pope points to three features of everyday AI use that deserve careful attention: "the ease with which results are obtained, the impression of objectivity and the simulation of human communication." Speed and ease are real benefits, "yet they can also encourage excessive reliance and the search for ready-made answers, and weaken personal creativity and judgment."

He also distinguishes information from real knowledge. Many teachers, he notes, already see students who "know many things" but struggle to find direction in their lives, because "an incessant flow of information replaces the essential exercise of research, reflection and discernment." Wisdom, he says, needs "rhythms that incorporate silence, in-depth study, reading and judicious analysis."

For this reason, Catholics should be cautious about outsourcing moral judgment or spiritual formation to machines. Questions involving conscience, doctrine, virtue, repentance, and salvation belong properly to the realm of the human soul under God's grace.

The Pope puts this in a way every student should remember. Learning to use AI well is not only learning *how* to use it:

> "Educating people about the use of AI, then, involves teaching them to decide when and for what purpose it ought not to be used."
>
> — *Magnifica Humanitas*, no. 140

He warns that "the speed and ease with which answers or summaries can be obtained risk extinguishing the desire to ask questions," and asks us to protect young people "from the promise of the perfect machine, from that subtle temptation which renders human thought seemingly superfluous precisely when it is most needed." AI may assist learning or organization, but it cannot replace the authority of Sacred Scripture, Sacred Tradition, the Magisterium of the Church, prayer, or authentic human wisdom.

Another important concern is relativism. AI systems are often designed to avoid conflict or present multiple viewpoints neutrally. As a result, they may frame radically opposed moral systems as interchangeable "perspectives." Catholic teaching rejects the idea that truth is merely subjective preference. Contradictions are real, and error has consequences.

Magnifica Humanitas goes further and names what happens when people with great technological power stop serving truth. Those who control powerful technologies and media can influence "a significant number of people concerning the truth about humanity, the world, the meaning of existence, the family and even God."

> "This is pure power detached from truth, which subtly or overtly imposes what it wishes others to accept as true."
>
> — *Magnifica Humanitas*, no. 133

At the root of this, the Pope says, is the modern belief that we can "construct reality, and that whatever best suits their claims corresponds to what is true." When that happens, "universally valid truths, which precede us and which conscience must accept, are no longer recognized." He warns that this is not only a private problem: "indifference to the truth leads, slowly but surely, to a descent into totalitarianism." An AI may present Catholic morality, utilitarian ethics, nihilism, and pagan spirituality side-by-side with equal detachment, but a Catholic must remember that truth is not determined by balanced presentation or social tolerance.

Catholic philosophy also teaches that truth is connected to being itself. Everything that exists participates in truth because creation reflects the intelligibility of God. Yet AI does not possess personal being in the way human persons do. It has no immortal soul, no eternal destiny, no capacity for grace, and no vocation to communion with God. It is an artifact created by human ingenuity, not a person created in the image of God. It may imitate aspects of intelligence while lacking the deepest dimensions of what intelligence truly is.

None of this means AI is useless or inherently evil. Like many technologies, it can serve legitimate human purposes when used prudently. AI may assist research, education, communication, accessibility, and organization. Catholics may make use of such tools responsibly while remaining aware of their limitations. The key distinction is that AI can communicate truths instrumentally, but it can never become the source or measure of truth itself.

This is why Pope Leo XIV insists that "it will always be human intelligence, with its conscience and freedom, that guides technical innovations and responsibly determines their use and limits."

Ultimately, Catholicism teaches that truth is not fulfilled in computation or data processing, but in communion with God. Information is not wisdom, simulation is not personhood, and intelligence alone is not holiness. The fullness of truth is found not in machines, but in the eternal Logos, Jesus Christ, through whom all things were made and by whom all things are judged.

Magnifica Humanitas begins and ends here. It opens by recalling that it is "only in the mystery of the Word made flesh that the mystery of humanity truly becomes clear." And near its close, the Pope gives this charge:

> "Let us remain faithful to the truth! Living amid incessant flows of information, opinions and images, we know how easy it can be to influence decisions and preferences through increasingly sophisticated algorithms. In this context, it is

> imperative to cultivate hearts that love the truth, prefer what is right despite the most appealing content and pursue wisdom rather than immediate results."
>
> — *Magnifica Humanitas*, no. 237

5. Bias

Bias in AI is not an accidental flaw that can be completely eliminated; it is an unavoidable consequence of the fact that AI is built by fallen human beings and trained on human culture. Every AI system reflects the assumptions, priorities, and worldview of the people and institutions that design it. Since humanity after the Fall is wounded by pride, selfishness, ideological confusion, and moral disorder, those distortions inevitably influence the systems we create. AI therefore does not emerge from some realm of pure neutrality. It reflects the civilization that formed it.

This is the heart of Pope Leo XIV's teaching that technology "is never neutral, because it takes on the characteristics of those who devise, finance, regulate and use it." Notice the four groups he names. Bias can enter an AI system at each stage: when engineers **design** it, when investors decide what to **fund**, when lawmakers decide what to **allow**, and when ordinary people decide how to **use** it.

Catholic thought has long recognized that no intellect approaches reality without formation. Human beings are shaped by family, education, philosophy, religion, political structures, and moral habits. AI systems are similarly “formed” by training data, corporate incentives, moderation policies, cultural norms, and the judgments of developers about what should or should not be allowed. Even deciding what information counts as trustworthy, what speech is harmful, or what moral perspectives are acceptable already introduces bias. The claim that a system is entirely neutral often conceals its deepest assumptions.

The Pope warns that hidden bias is most dangerous exactly when a system *looks* objective. An AI's calm, confident answers can make us forget that "they reflect the cultural assumptions of those who designed and trained them, with all their strengths and limitations."

> "There is also a subtler danger, for when AI systems present themselves as neutral and objective, they end up reflecting and reinforcing the stereotypes or ideological bias of their designers and developers."
>
> — *Magnifica Humanitas*, no. 102

He names where this matters most: "important and sensitive decisions — concerning employment, credit, access to public services or even a person’s reputation." When an AI system screens job applications, approves loans, or flags people for review, a hidden bias is

no longer just an inaccurate sentence. It becomes an injustice done to a real person. This is why the encyclical treats biased AI as a matter of **social justice**, warning against "social groups penalized by opaque algorithms that perpetuate prejudice and discrimination."

At the same time, Catholicism distinguishes between unjust prejudice and rightful moral judgment. Modern discussions sometimes treat any preference or discrimination between ideas as inherently oppressive. But the Church insists that some distinctions are morally necessary. To prefer truth over falsehood, virtue over vice, justice over injustice, or life over death is itself a kind of "bias," but it is a righteous one. Catholicism unapologetically favors what is objectively good and true. Therefore an AI that presents all moral claims as equally valid is not truly neutral; it is already biased toward relativism and indifferentism.

Much contemporary AI development takes place within highly secular, technocratic, and post-Christian cultural environments. As a result, many systems naturally reflect the assumptions dominant within those environments. This can especially affect discussions surrounding abortion, sexuality, gender ideology, religious exclusivity, sin, authority, and the nature of the human person. A Catholic should therefore expect tension between traditional Christian teaching and many prevailing cultural narratives embedded within technological systems. This does not mean every AI output is false or malicious, but it does mean discernment is necessary. The machine often mirrors the moral and philosophical condition of the society producing it.

Magnifica Humanitas points to a deeper problem than which values happen to be popular. A small number of private companies now control most powerful AI, with resources "that surpass those of many Governments." Those who control digital platforms "have a considerable ability to affect the collective imagination and to present a particular vision of reality as desirable." Companies try to "align" their AI with human values, but the Pope asks: whose values, and who decides? If a few people make those choices out of public view, their moral vision "will become the invisible infrastructure of these systems."

> "A more moral AI is not enough if that morality is determined by a few."
>
> — *Magnifica Humanitas*, no. 107

Bias can also operate through omission rather than direct falsehood. Sometimes what is left unsaid is as influential as what is spoken. AI may avoid moral clarity, soften distinctions between truth and error, or frame Catholic doctrine as merely one opinion among countless competing perspectives. Over time, this can subtly shape users toward relativism, where objective truth is replaced by personal preference or social consensus. Catholic teaching rejects this approach because truth is not created by popularity or cultural fashion. Truth exists independently of human approval because it is grounded ultimately in God.

At the same time, Catholics must also guard against their own biases. It would be a mistake to respond to AI with paranoia, tribalism, or ideological obsession. Human beings themselves are prone to anger, fear, partisanship, and intellectual pride. The Christian pursuit

of truth requires humility, discipline, charity, and a willingness to conform oneself to reality rather than merely defending one's preferred narrative. Discernment must therefore be exercised both toward technology and toward our own hearts.

Pope Leo XIV models this humility. He writes that the Church "does not claim to possess a monopoly on truth," because truth "is a gift to be shared, not a possession to be monopolized." This does not make truth relative. It means that whoever holds the truth should hold it as a servant, not as a weapon, and should be the first to examine his own blind spots.

A particularly important difference between humans and AI is that human beings can repent and be transformed by grace. A sinner can convert, grow in wisdom, and become holy. The Christian tradition is full of examples of dramatic conversion, such as Saint Paul the Apostle, who went from persecuting Christians to becoming one of the greatest apostles. AI, however, has no soul, conscience, or moral interiority. It cannot seek virtue or repent of error. Its "alignment" is always externally imposed by those who control and shape it.

The Pope draws out a striking consequence of this difference:

> "For an algorithm, an error is a flaw to be corrected; for a person, however, an error can be a catalyst for profound change. A person's future is not calculable, but depends on one's freedom — elevated by the inexhaustible grace of God — and on the relationships cultivated."
>
> — *Magnifica Humanitas*, no. 128

An AI system that sorts people by their past data can never see the person Saint Paul would become. Only God's grace and human freedom can.

For that reason, Catholics should approach AI as a tool rather than an authority. It can be useful for research, organization, communication, and learning, but it should never replace Scripture, Sacred Tradition, the teaching authority of the Church, sound philosophy, or the formation of conscience. A faithful Catholic approach recognizes that eloquence is not wisdom, consensus is not truth, and technological sophistication is not holiness. Ultimately, every human system — including AI — must be judged according to reality as established by God, who alone is the fullness of Truth.

6. Ethics

From a Catholic perspective, the ethical questions surrounding AI are fundamentally questions about the human person. Technology itself is not morally good or evil in the abstract; its moral character depends on the ends it serves and the way it shapes human life. The Church has always taught that tools should remain subordinate to the dignity of the human person, never the reverse. AI therefore becomes dangerous when efficiency, control,

profit, or convenience begin to outweigh truth, charity, justice, and the sanctity of human life.

One of the first ethical concerns is the temptation to treat AI as morally authoritative. Because AI can speak fluently, synthesize information rapidly, and appear objective, people may begin to defer to it in matters requiring conscience and prudence. Yet Catholic teaching insists that conscience is not a machine process but an act of the human soul informed by reason, natural law, and divine revelation. A machine cannot bear moral responsibility, cannot sin, cannot repent, and cannot love. Therefore it cannot replace authentic moral judgment. A Catholic should resist the subtle cultural drift toward allowing algorithms to determine what is true, acceptable, or just.

Another major ethical issue is the danger of depersonalization. Human beings are not data points or economic units; they are persons made in the image of God. AI systems can encourage societies to reduce people to profiles, behavioral predictions, productivity metrics, or ideological categories. This can erode compassion and undermine the Christian understanding of the person. In workplaces, education, medicine, and even family life, there is a risk that human relationships become mediated through systems optimized for efficiency rather than charity. Catholic thought warns against any social order that sacrifices human dignity to technocratic management.

Pope Leo XIV, building on Pope Francis, calls this the **technocratic paradigm**: "the tendency to let the logic of efficiency, control and profit alone shape personal, social and economic decisions." When technology "becomes the standard by which everything is judged, it begins to dictate what matters and what can be discarded," reducing "human beings to mere cogs in a system driven toward ever greater efficiency." He calls the temptation behind it the **"Babel syndrome"**: "the pretense that a single language — even a digital one — can translate everything, including the mystery of the person, into data and performance."

There is also the ethical problem of manipulation. AI can shape perception, influence emotions, filter information, and personalize persuasion at enormous scale. In the wrong hands, this becomes a powerful instrument for propaganda, surveillance, consumer exploitation, or political coercion. Catholic moral teaching has long defended the freedom of the human person and the integrity of the conscience. Systems designed to manipulate populations psychologically or ideologically violate that dignity, even if done under the language of safety, progress, or social harmony. A society that increasingly governs through invisible algorithmic influence risks weakening moral responsibility and authentic freedom.

The ethics of truth are especially important. AI systems can generate convincing falsehoods, distortions, and fabrications with great speed. Catholicism teaches that truthfulness is not optional because lying damages both the speaker and the community. Technologies that normalize deception, impersonation, or manufactured reality can corrode trust itself. This is spiritually serious because human society depends upon confidence that words correspond to reality. If people become unable to distinguish truth from fabrication, public life

deteriorates into manipulation and suspicion. In that sense, AI intensifies an already existing moral crisis within modern culture: the weakening of commitment to objective truth.

There is also an ethical danger in AI that *imitates* human care. Chatbots can offer "words of advice, empathy, friendship and even love." The Pope grants this "can be engaging and at times genuinely helpful," but "when words are simulated, they do not build genuine relationships, but only their appearance."

> "Here, the danger is not so much that a person may believe they are communicating with another person, but rather that they may gradually lose the very desire to form genuine human connections."
>
> — *Magnifica Humanitas*, no. 100

Catholic ethics would also raise concern about pride and the illusion of self-sufficiency. There is a recurring temptation in human history to believe that technological power can solve the deepest problems of existence. Modern rhetoric around AI sometimes approaches a kind of secular salvation narrative, promising liberation from limitations, suffering, labor, or even death itself. Christianity sees this as spiritually dangerous because salvation cannot come from technique. Sin, mortality, and the need for redemption are not engineering problems. When societies begin to place ultimate hope in technological systems, they risk creating a counterfeit religion detached from God.

Magnifica Humanitas identifies the ideas behind this counterfeit religion: **transhumanism**, which seeks to "enhance" human beings through technology, and **posthumanism**, which imagines humanity surpassing itself in "a new evolutionary stage." The Pope says the key issue "is not the use of technology as such, but the vision that underlies it."

> "If the human being is treated as something to be perfected or surpassed, it becomes easier to accept that some lives are less useful, less desirable or less worthy."
>
> — *Magnifica Humanitas*, no. 117

Against this, the Pope teaches that "humanity flourishes not despite limitations, but often through them," and that building the common good "means accepting the limits and weakness of humanity without considering them an error to be corrected." The true "more than human" is not found in technology but in grace: "what saves humanity is not enhanced self-sufficiency, but a relationship that liberates, a communion that transforms." He also quotes Romano Guardini: "Contemporary man has not been trained to use power well."

The encyclical adds an ethical cost that is easy to overlook: the environment. Current AI systems "require enormous amounts of energy and water," and as large language models grow, so do the "data centers and energy-intensive infrastructure" they depend on. Caring for our common home is part of the ethics of AI.

At the same time, Catholic teaching does not require fear of technology. AI can be used for genuinely good purposes: advancing medical research, improving accessibility for the disabled, assisting education, organizing knowledge, or helping people communicate more effectively. The moral question is always whether technology serves authentic human flourishing in accordance with natural law and divine truth.

The Pope offers one question for evaluating any technology, including AI: "Do they truly help individuals and peoples to become more humane and fraternal, while respecting our common home and future generations?" The Church has historically embraced legitimate advances while warning against the misuse of power and the forgetting of God.

Ultimately, Catholic ethics insists that no machine can replace the moral and spiritual vocation of the human person. Human beings are called not merely to process information, but to grow in wisdom, virtue, holiness, and communion with God. AI may imitate aspects of intelligence, but it cannot love, sacrifice, worship, or attain sanctity. Therefore the ethical boundary is clear: technology must remain a servant of humanity and never become an object of dependence, domination, or reverence.

7. Morality

Morality is inseparable from truth, human nature, and God's eternal law. Because AI is not a person and does not possess a rational soul, it cannot itself be moral or immoral in the proper sense. It has no conscience, no free will, no capacity for virtue, and no ability to sin. A machine cannot choose the good out of love for God, nor can it rebel against God through deliberate evil. Moral responsibility always belongs ultimately to the human beings who design, direct, deploy, and use the technology.

This distinction is essential because modern culture increasingly speaks about AI as though it were an independent moral actor. In reality, AI does not "decide" morally; it processes instructions and patterns according to objectives established by human beings. Even when an AI appears to reason ethically, it is not exercising conscience. It is simulating moral language based on data and programming. Catholic teaching understands conscience as an act of the human intellect applying moral truth to concrete situations under the guidance of reason and divine law. Since AI lacks intellect and soul, it cannot possess conscience in the true sense.

Magnifica Humanitas is clear: AI systems do not "have a moral conscience, since they do not judge good and evil, grasp the ultimate meaning of situations, or bear responsibility for consequences." Near its close, the encyclical says it even more simply:

> "No computational system, however sophisticated, can create a heart that gives itself, or a conscience that discerns good from evil."
>
> — *Magnifica Humanitas*, no. 233

(You will study the question of who is responsible for AI's decisions in much more depth in a later class.)

Catholic morality begins with the belief that there is an objective moral order rooted in God Himself. Right and wrong are not created by consensus, efficiency, emotional preference, or social utility. Human beings discover moral truth through natural law, revelation, reason, and the teaching authority of the Church. AI systems, however, are built within cultural and institutional frameworks that may or may not respect objective morality. As a result, they can reflect distorted moral assumptions present within society, including relativism, utilitarianism, materialism, or ideologies hostile to Christian anthropology.

One serious moral danger is that AI can encourage people to outsource conscience and personal responsibility. If individuals begin relying on machines to determine what is right or wrong, they risk weakening the moral faculties that God intended human beings to cultivate through prayer, discipline, study, and virtue. Catholic tradition teaches that moral maturity requires formation of conscience, self-denial, wisdom, and accountability before God. A machine cannot replace that process. Delegating deeply moral decisions entirely to algorithms risks treating morality as a technical calculation rather than a matter of truth and virtue.

Catholic thought would also warn against the reduction of human beings to data or utility. Modern technological systems often evaluate people according to efficiency, productivity, behavioral prediction, or ideological conformity. But the Church insists that every person possesses intrinsic dignity because he or she is made in the image of God. Human value does not come from usefulness, intelligence, social status, or economic contribution. AI systems that dehumanize, manipulate, surveil excessively, or categorize persons merely as objects of management can undermine this fundamental moral truth. It is possible to lose your dignity to AI.

Pope Leo XIV calls one modern ideology "particularly insidious": the idea "that every person must earn or justify his or her own worth, to the point of attributing greater value to those who are more efficient or effective." Efficiency is exactly what AI is built to measure and optimize, which is why this temptation is so strong in a world shaped by algorithms.

A careful distinction from the encyclical helps here. The Pope describes several kinds of dignity. *Moral* dignity (how we direct our choices), *social* dignity (how society treats us), and *existential* dignity (how we perceive our own worth) "can be enhanced or diminished." But **ontological dignity**, the dignity that belongs to every human being "simply by virtue of existing, of having been willed, created and loved by God," can never be lost. "No sin, failure, humiliation or exclusion can diminish" it. A system can fail to respect your dignity, and you can damage your moral dignity by how you choose to act, but no machine can take away the dignity God gave you.

> "The value of persons, however, does not depend on what they achieve or produce."

> — *Magnifica Humanitas*, no. 51

The morality of truthfulness is especially relevant in the age of AI. Technologies capable of generating convincing falsehoods, impersonations, propaganda, or fabricated realities create profound ethical concerns. Catholic morality condemns lying because it damages trust and obscures reality. Human society depends upon confidence that words correspond to truth. When technology normalizes deception or blurs the distinction between reality and fabrication, it contributes to moral confusion and social fragmentation. This becomes especially dangerous when falsehood is used to manipulate consciences, influence elections, distort history, or exploit vulnerable people.

Another moral concern involves pride and the temptation toward self-sufficiency. Throughout history, humanity has repeatedly sought salvation through power, empire, wealth, or knowledge detached from God. Contemporary rhetoric surrounding AI sometimes approaches a kind of technological messianism, promising liberation from suffering, labor, limitation, or even mortality itself. Catholicism sees this as spiritually dangerous because no technological achievement can redeem humanity from sin or replace the need for grace. Human flourishing ultimately depends not on mastering creation absolutely, but on rightly ordering ourselves toward God.

At the same time, Catholic morality does not reject technology outright. The Church has historically affirmed the legitimate use of human ingenuity when directed toward authentic good. AI can assist medicine, education, accessibility, communication, scientific research, and many practical tasks that genuinely benefit society. The moral question is always whether technology serves the true good of the human person and respects the natural moral order established by God. A tool becomes dangerous when efficiency or profit is elevated above truth, charity, justice, and human dignity.

Catholic morality also emphasizes that virtue cannot be automated. Holiness involves love, sacrifice, humility, repentance, courage, and obedience to God. These are realities of the soul, not functions of computation.

The Pope makes the same point when he says AI does not "know from within what love, work, friendship or responsibility mean." Virtue grows through a lifetime of choices, mistakes, forgiveness, and fidelity. A machine can repeat the vocabulary of virtue without ever possessing it. The encyclical even suggests that limits are where virtue is born: "It is precisely within our limitations that the following find a place: compassion, as well as a sincere concern for the needs of others." AI may imitate aspects of reasoning or communication, but it cannot become virtuous because it cannot love. It cannot worship, pray, suffer redemptively, or grow in sanctity. The moral life remains fundamentally personal and spiritual.

Ultimately, the Catholic understanding of morality insists that technology must remain subordinate to the human person and the human person subordinate to God. AI should be treated as a tool to assist human activity, not as a moral authority or substitute for wisdom.

The measure of whether a technology is used rightly is whether it leads people toward truth, virtue, justice, and authentic human flourishing. Any system that weakens conscience, obscures truth, diminishes human dignity, or encourages separation from God becomes morally dangerous, regardless of how advanced or useful it may appear.

No LLM knows whether its response is truthful or not, moral or not or ethical or not. They do not possess understanding, beliefs, or consciousness. Because training data includes human-written content, LLMs can reproduce biases, stereotypes, or harmful language present in that data. They may also be manipulated by cleverly phrased inputs to produce unsafe outputs. Unless a model's training data is built or cleansed through the Magisterium of the Catholic Church, no LLM will ever create or imitate the Word.

Notice what *Magnifica Humanitas* adds. Better training data and more careful "alignment" are worth pursuing, but they are not enough on their own. The Pope says we "cannot be satisfied with merely calling for the moralization of machines," because whoever trains a model still decides which morality it reflects. That is why human beings, not machines, must remain the ones who judge, decide, and answer for how these tools are used.

LLMs are technologies to be stewarded: they reflect human ingenuity and can promote human flourishing when used ethically. Moral discernment should guide how educators, developers, and students apply these tools—ensuring they support truth, human dignity, and the common good rather than undermine them.

> "In the era of artificial intelligence, when human dignity is threatened by new forms of dehumanization, ours is the pressing duty to remain profoundly human."
>
> — *Magnifica Humanitas*, no. 15

8. What This Means for You as a Student

Magnifica Humanitas speaks directly about schools and young people. The Pope writes that school "is the place where new generations can learn to seek and love the truth, to reflect on the meaning of life and to recognize the dignity of every person." He also gives a surprising instruction:

> "Schools are not called to follow the pace of the digital world, but to offer that which the digital sphere by itself cannot provide, namely a shared time for learning and developing trustworthy relationships."
>
> — *Magnifica Humanitas*, no. 147

Here is how the encyclical's teaching applies to the way you use AI in this course and beyond:

- **Think first, then ask.** The Pope warns that easy answers can extinguish "the desire to ask questions." Write your own first draft before you prompt an AI.
- **Verify.** Truth "requires verification, cross-checking of sources and responsible argumentation." Treat every AI answer as a claim to check, not a fact to copy.
- **Ask whose assumptions are built in.** An AI's "impression of objectivity" can hide the views of those who designed and trained it.
- **Know when not to use it.** Prayer, conscience, friendship, and the formation of your own judgment are not tasks to hand to a machine.
- **Keep your relationships human.** A chatbot can simulate friendship, but it cannot be a friend.
- **Recognize manipulation.** The Pope asks that young people learn "how to recognize manipulation, defend their dignity and respect that of others in digital environments."

9. Exit Ticket

1. According to Catholic teaching, why can AI never truly “know” truth the way a human person can?

2. Explain one way AI systems can reflect bias from society or their training data.

3. Why does the Church teach that morality cannot be fully delegated to machines or algorithms?

4. What is one ethical danger of treating AI as morally authoritative?

5. Explain the difference between information and wisdom from a Catholic perspective.

6. Pope Leo XIV says technology "is never neutral, because it takes on the characteristics of those who devise, finance, regulate and use it." Choose one of those four groups and explain how bias could enter an AI system at that stage.

7. Why does *Magnifica Humanitas* say that an AI system's "impression of objectivity" makes its bias harder to notice?

8. What does Pope Leo XIV mean when he calls truth "a common good and not the property of those with power or influence"?

9. According to the encyclical, why does learning to use AI well include learning "when and for what purpose it ought not to be used"?

10. Explain the Tower of Babel and Nehemiah images from *Magnifica Humanitas*. Which one describes how you want to use AI, and why?

10. Reflection

- Why does Catholic teaching insist that technology must remain a servant of humanity rather than become an object of dependence or reverence?
- Pope Leo XIV writes: "For an algorithm, an error is a flaw to be corrected; for a person, however, an error can be a catalyst for profound change." Think of a mistake that changed you for the better. Why could a machine never experience that kind of change?
- The encyclical says that "the construction of Babel or the rebuilding of Jerusalem begins within each one of us." What is one concrete habit in your own use of technology that builds Babel, and one that helps rebuild Jerusalem?

11. Vocabulary

1. **Bias** A tendency or inclination that influences judgment or behavior; in AI, bias often comes from the data or assumptions used to train a system.
2. **Conscience** The moral faculty of the human soul that helps a person discern right from wrong according to reason and God's law.
3. **Relativism** The belief that truth or morality is subjective and depends on personal opinion or cultural preference rather than objective reality.
4. **Human Dignity** The inherent worth and value of every human person because they are created in the image and likeness of God.
5. **Objective Truth** Truth that exists independently of personal feelings, opinions, or social consensus.
6. **Moral Agency** The ability to make moral decisions and be responsible for one's actions; humans possess this, but AI does not.
7. **Natural Law** The moral law written into human nature by God and discoverable through reason.
8. **Discernment** The process of carefully evaluating and judging situations according to truth, wisdom, and moral principles.
9. **Technocratic Paradigm** The tendency to let the logic of efficiency, control, and profit alone shape personal, social, and economic decisions, so that technology becomes the standard by which everything is judged.
10. **Magisterium** – The teaching authority of the Catholic Church.
11. **Intellectus** – The spiritual faculty of understanding and reasoning in classical Catholic philosophy.

12. **Utilitarianism** – A moral theory that judges actions mainly by usefulness or outcomes rather than objective moral truth.

13. **Anthropology (Christian Anthropology)** – The study of the human person from a Christian perspective.

14. **Virtue** – A habitual disposition toward doing good.

15. ***Magnifica Humanitas*** – "Magnificent Humanity"; Pope Leo XIV's first encyclical (May 15, 2026), on safeguarding the human person in the time of artificial intelligence.

16. **Encyclical** – A formal teaching letter from the Pope, addressed to the Church and often to all people of good will.

17. **Babel Syndrome** – Pope Leo XIV's term for the temptation to build without God: the idolatry of profit, a uniformity that erases differences, and the pretense that one digital language can translate everything, including the human person, into data.

18. **Common Good** – The sum of social conditions that allow people, as groups and as individuals, to reach their fulfillment; *Magnifica Humanitas* teaches that truth itself is a common good.

19. **Alignment** – The effort to make an AI system's behavior match human values; the encyclical asks whose values are chosen and who decides.

20. **Ontological Dignity** – The dignity every human being has simply by existing as a person willed, created, and loved by God; it can never be lost or taken away.

21. **Transhumanism** – A movement that seeks to "enhance" human beings through technology in order to overcome human limits.

22. **Disinformation** – False or misleading information spread deliberately; the encyclical warns that AI is "a powerful amplifier" of it.